

Working with NoSQL Databases Using Python



Pinal Dave

Data Expert

@pinaldave | blog.sqlauthority.com



Why Do We Need NoSQL?

Traditional relational databases struggle with scalability and flexibility

NoSQL is designed for high-volume, distributed, and unstructured data

Ideal for real-time applications, IoT, social media, and big data analytics




NoSQL vs. SQL - Key Differences

Feature	SQL (Relational DBs)	NoSQL (Non-relational DBs)
Schema	Fixed, predefined	Flexible, schema-less
Scalability	Vertical (scale-up)	Horizontal (scale-out)
Data storage	Tables (rows & columns)	JSON, key-value, columnar, graph
Transactions	ACID	BASE (eventual consistency)
Use cases	Structured data, financial apps	Big data, real-time apps



Scaling NoSQL Databases



**Large-scale
applications demand
distributed data
storage**




**NoSQL databases
scale horizontally
(adding more nodes)**



**Proper indexing and
partitioning are
crucial for
performance**

Balancing Reads and Writes

Depending on the workload, different strategies can enhance performance and scalability while maintaining data consistency.




Read-heavy workloads

**Use caching
(Redis, Memcached)**



Write-heavy workloads

**Use batch writes and
optimize schema**



Eventual consistency

**Improves performance
but may lead to
stale data**

